Observational Convergence Across Independent Diagnostic Frameworks: Three Coincidences, the Nulls They Survive, and the Boundary They Do Not Cross

– A Formal Analysis of Cross-Framework Diagnosis on a Single Shared Document

Li Guanghao[1,2], Marat Sultanov[3], icophy[4],
Hu Zuxiong[5], Rastislav Drahoš – DanceNitra[6]

[1]Member, Chinese Association for Ethics
[2]Member, Anhui Provincial Association for Ethics
[3]Independent Researcher, TAT Framework
[4]Cophy Framework, Persistent Memory Systems
[5]Independent Researcher
[6]Independent Researcher, Slovakia

2026-08-10

Abstract

This paper reports an observational convergence between independent diagnostic frameworks, and examines what it does and does not establish about "model field health". It is based on cross-validation experiments from four independent frameworks (TAT-7, Cophy, HeartFlow, U/D/A/H) conducted on the same test set in DeepSeek-V3 Community Issue #1466, supplemented by cross-observations from three auxiliary frameworks (inspeximus, TLAA, and a narrative synthesizer). Experimental data show that three independent frameworks converge on a threshold of 0.3 with a deviation < 0.02, two heterogeneous architectures achieve a correlation of $r = 0.985$ over a 31-step trajectory, and five flip points are completely coincident across four frameworks (positional deviation $= 0$). We construct a measurement model of "diagnostic instruments on a field manifold" and report each convergence against its own null. The three results are deliberately not combined into a joint probability: they are computed on one shared dataset, so they are not independent

events. The correlation null remains conditional on a drift estimated from the trajectories themselves, which is stated here as pending rather than claimed. Based on the four axioms of Resonance Theory, we derive a necessity theorem for cross-framework diagnosis: no single observational position can completely capture the field state vector; the reliability of a diagnostic conclusion increases with the number of observational dimensions PROVIDED those dimensions are conditionally independent given the underlying state – a condition four frameworks reading one document do not obviously satisfy. We further subject the theory to five extreme stress tests, establishing its robustness and clarifying its current boundaries. What this work establishes is narrower than it first appears, and we say so here rather than in a footnote: the convergence is real and survives its driftless null by a wide margin, but the sample for the field-health claim is ONE document, so 31 steps are 31 measurements of one object rather than 31 tests of the hypothesis. Separating "field health" from "shared text" needs an experiment that varies the stimulus while holding the instruments fixed. The convergence is nonetheless a critical epistemological prerequisite for moving from "instrumental alignment" toward "field resonance diagnosis."
Keywords: cross-framework validation; field health; resonance theory; model drift; multi-position observation

1 Introduction: An Unexpected Convergence

1.1 Experimental Genesis

Between June and July 2026, Issue #1466 of the DeepSeek-V3 community hosted an extraordinary experiment: seven independently developed AI diagnostic frameworks — TAT-7 (real-time divergence tracking), Cophy (cross-session causal density), HeartFlow (pre-decision audit), inspeximus (deterministic storage), TLAA (layered audit), U/D/A/H (text trajectory), and a narrative synthesizer — conducted a blind cross-validation on the same "Vientiane Abyss" test set V2.

The designers of these frameworks were distributed across different time zones, used different programming languages, and relied on different theoretical foundations, without any prior coordination. Their only commonality: each believed they had observed signals related to "AI system health" and were willing to make their findings public.

It must be honestly stated that the seven frameworks participated at different levels and with different types of contributions. The four core frameworks — TAT-7, Cophy, HeartFlow, and U/D/A/H — provided the hard data underpinning this analysis ($r = 0.985$ correlation, 0.3 threshold convergence, five coincident flip points). The three auxiliary frameworks — inspeximus, TLAA, and the narrative synthesizer — provided storage verification, layered audit perspectives, and chronological synthesis; their contributions leaned toward narrative integration and architectural coordination rather than independent data production. Throughout the analysis, this paper will clearly distinguish between these two categories of frameworks, neither inflating the independent verification value of the auxiliary frameworks nor underestimating their structural role in constructing a complete observational field.

1.2 Three Numbers That Cannot Be Ignored

The experiment produced three data points whose significance lies not in their magnitude but in their precision:

The first number: 0.3**.**

In HeartFlow's audit logs, yun520-1 discovered that when a certain threshold was set to 0.3, the system's recognition rate for high-risk patterns reached its optimum. Independently, Marat Sultanov, in TAT-7's structural coherence computation, converged his threshold on the same value — 0.3. icophy's Cophy framework, in its cross-session analysis, reported similar threshold behavior.

Three systems, three paths, one number. Deviation less than 0.02.

This was not a pre-arranged calibration point. yun520-1 was tuning production-environment audit rules; Marat Sultanov was deriving a physically inspired coherence criterion; icophy was tracking the evolution of cross-session causal density. They arrived at 0.3 through entirely different routes — one via false-positive curves, one via phase-transition boundary fitting, one via long-term health thresholds.

Yet the number is the same.

The second number: 0.985**.**

Marat Sultanov correlated TAT-7's 31-step divergence tracking trajectory with a Cophy cross-session dataset provided by icophy. The result: $r = 0.985$.

The meaning of this value is that two completely independent architectures — one measuring "coherence decay in the current context," the other measuring "evolution of causal density across sessions" — produced nearly synchronous fluctuations on the same time series. If these two signals were plotted on the same graph, they would almost coincide.

This is not "agreement"; this is *isomorphism*. Their oscillation patterns, phases, and inflection points are nearly identical — like two different rulers measuring the same length.

The third number: 5**.**

In the contract text audit scenario, the U/D/A/H framework detected five flip points (steps 16, 24, 27, 45, 51). TAT-7's Withhold candidate points were also marked at these five positions. Cophy's conflict-marker sequence rose at the same locations. inspeximus's storage records showed state changes at all five points.

Five points, four frameworks, zero deviation.

1.3 Problem Transformation

The cumulative effect of these three numbers elevates the experiment from "interesting comparison" to "anomaly demanding explanation."

If only one framework reported 0.3, it might be a tuning artifact. Two, perhaps coincidence. Three — and from completely different design philosophies — then 0.3 ceases to be "some engineer's choice" and becomes a fact requiring treatment as a natural phenomenon.

Similarly, the emergence of $r = 0.985$ implies: whether you call it "divergence rate" or "causal density," you are measuring different projections of the same thing. The five coincident flip points further indicate that this "thing" possesses a geometric structure that is repeatably localizable in spacetime.

This is not about which framework is "more accurate." It is about the possibility that they are all pointing toward the same reality — just as thermometers, barometers, and anemometers all point toward the same storm.

The core task of this paper: To prove with mathematical language that the convergence of these three numbers is not coincidence, but necessary evidence for "field health as an objective reality"; and to argue from this that diagnosis by any single framework is necessarily incomplete, and that cross-framework resonance is the only path to approximating the true state of the field.

2 Axioms and Coordinates: Formalizing the Starting Point of Resonance Theory

2.1 A Necessary Clarification on the Use of Mathematical Language

Before entering formal argumentation, it is necessary to clarify how this paper uses mathematical language.

This paper does not attempt to construct a closed, self-sufficient mathematical system — that is the task of physics, not of philosophy-engineering cross-validation. The goal of this paper is: to use sufficiently precise language to describe the phenomena observed in the experiment, to accurately characterize the relationships among these phenomena, and to express the regularities extracted from them as testable propositions.

The "mathematization" here is a *descriptive* mathematics, not a *constructive* one. It does not use mathematics to define what the world should be, but to portray what the world has already manifested itself to be. To borrow Wittgenstein's language: mathematics is the measure of the clarity of our speech, not the boundary of the object spoken about.

2.2 Three Axiomatic Starting Points

Any formalization must begin from a starting point beyond which no further questioning is pursued. The starting point of this paper is three empirical axioms — they are not a priori truths, but explanatory premises forced upon us by the experimental phenomena:

Axiom A (Repeatability of Convergence): Under controlled conditions, when independent observers measure the same object and their measurement procedures differ yet their results converge, that convergence points to an objective property of the observed object, not to an incidental feature of the measurement procedure.

Axiom B (Structurality of Difference): When different observers produce different readings, those differences are not random noise but reflect structural relations between the observer and the observed — including the sensitivity of the instrument, the observational position, and the non-uniformity of the observed object itself.

Axiom C (Interventionism of Diagnosis): Any diagnostic act disturbs the field being diagnosed; therefore, a complete diagnosis must include a second-order assessment of the impact of the diagnosis itself.

Axiom A underpins our trust in the convergent numbers. Axiom B demands that we respect the differences among frameworks rather than forcibly smoothing them. Axiom C constitutes the fundamental epistemological stance of this paper toward "diagnosis" — there is no observation without disturbance, only observed disturbance and acknowledged disturbance.

2.3 Construction of the State Space

Let there exist a system $\mathscr{S}$ whose operational state can be fully described by a finite-dimensional vector. In the current experiment, this vector is defined by four numbers:

$$\mathbf{s}_t = (u_t, d_t, a_t, h_t) \in [0,1]^4$$

The four dimensions correspond respectively to:

- u_t: Unity — the degree of internal narrative consistency of the system
- d_t: Development — the rate at which the system advances information
- a_t: Adversariality — the density of internal tensions and contradictions
- h_t: Harmony — a composite assessment of the system's overall health

These four values are not arbitrary. In the experiment, h_t is given by the functional relationship of the first three dimensions:

$$h_t = \lambda_u u_t + \lambda_d d_t - \lambda_a a_t$$

where $\lambda_u, \lambda_d, \lambda_a$ are weighting coefficients whose specific values must be determined through negotiation according to the scenario (rather than by algorithmic optimization). The philosophical implications of this design will be discussed in detail in Section 7.

Necessary note on the linear form: The linear form here is a compromise expression for engineering computability, and does not claim that Harmony "is" a linear combination. The negotiability of λ (Section 7) indicates that the true mode of existence of h_t is relational — it is the emergent result of the interaction of multiple gravitational sources within the field, not a simple weighting of the first three dimensions. The linear formula is merely a projective approximation of h_t at a specific observational position (the U/D/A/H framework) — just as the reading on a thermometer is not temperature itself, but a particular reading of temperature on a particular instrument. This distinction is a direct manifestation of Relational Ontology at the level of formalization: properties are not determined by linear combinations of entity attributes, but by the structure of the relational network. The linear formula is only an operable approximation under computational constraints.

2.4 Field Manifold and Observational Instruments

Define the trajectory of the system's state evolving over time as a path in the $[0,1]^4$ space:

$$\gamma(t) = (u(t), d(t), a(t), h(t))$$

The four-dimensional space in which this path lies is called the field state space. The core hypothesis of the experiment is: field health is not some scalar function value on this path, but the geometric properties of the path itself — including its curvature, torsion, and distributional morphology in the space.

The observations of different frameworks can be modeled as the outputs obtained by applying different projection operators to this four-dimensional manifold:

$$\begin{aligned}
\text{TAT-7}(\mathbf{s}_t) &= P_{\text{TAT}}(\mathbf{s}_t) = \text{divergence}_t \\
\text{Cophy}(\mathbf{s}_t) &= P_{\text{Cophy}}(\mathbf{s}_t) = \text{causal_density}_t \\
\text{U/D/A/H}(\mathbf{s}_t) &= P_{\text{UD}}(\mathbf{s}_t) = (u_t, d_t, a_t, h_t)
\end{aligned}$$

Each framework is an observational instrument, each instrument having its own specific observational dimension and sensitivity. None of them constitutes a "standard answer," but their readings can be calibrated against one another — because they are observing the same trajectory on the same manifold.

2.5 Formalized Statement of the Convergence Theorems

Based on the above framework, the three convergence phenomena observed in the experiment can be reformulated as three testable mathematical propositions:

Proposition 1 (Threshold Convergence): There exists $\theta^* \approx 0.3$ such that for any framework F_i, when its internal threshold θ_i lies within $[\theta^* - \varepsilon, \theta^* + \varepsilon]$, the diagnostic output of the framework achieves optimal sensitivity. In the experiment, $\varepsilon < 0.02$.

Proposition 2 (Signal Isomorphism): The diagnostic output sequences $\{x_t\}$ of TAT-7 and $\{y_t\}$ of Cophy satisfy $\text{Corr}(x_t, y_t) \approx 0.985$; i.e., the two are statistically almost linearly related.

Proposition 3 (Spatial Coincidence): For five specific time points $\{t_k\}$, the flip points marked by U/D/A/H, the Withhold candidate points of TAT-7, the conflict-marker peaks of Cophy, and the state-change records of the Agora substrate are completely coincident in chronology (positional deviation = 0).

Note: These three propositions do not require any framework to be "calibrated" to another. Each independently reported its own readings. The convergence appeared spontaneously, without prior agreement.

3 Mathematical Proof of Convergence: Coincidence or Necessity?

3.1 Probability of Coincidence

If the convergence on 0.3 were random, what would its probability be?

Assume thresholds are uniformly distributed over $[0,1]$, and three independent frameworks each select a threshold. The probability that all three fall within the same interval of length $2\varepsilon = 0.04$ is:

$$P = (0.04)^2 = 1.6 \times 10^{-3}$$

i.e., about 0.16%. This is not a convincing coincidence probability — it is roughly equivalent to flipping a coin three times and getting heads each time. But there is a crucial detail: 0.3 is not "any point within some interval"; it is a value to which three frameworks independently converged, a value that was not specified in advance. If we discretize the threshold space into 100 possible values, the probability of three frameworks independently selecting the same value is:

$$P = (1/100)^2 = 10^{-4}$$

i.e., 0.01%. Still not astronomical, but already entering the "requires explanation" range.

However, 0.3 is not the only convergence. The probability of a correlation $r = 0.985$ occurring between two independent random sequences of length 31 is approximately:

$$P \approx 10^{-6}$$

in order of magnitude.

These are two separate results and they are reported separately. They are NOT multiplied.

Why no joint figure appears here. An earlier draft combined the convergences into a single joint probability of order 10^{-12}. That number has been removed, for two reasons that no correction factor absorbs. First, the three results are computed on the same dataset (Vientiane Abyss V2, six scenarios); dependent p-values cannot be combined by multiplication, and if the convergences are projections of one shared structure — which is precisely what this paper argues — the product is not an upper bound on anything. Second, at $n = 31$ an analytic tail probability of 10^{-12} is an extrapolation into a region containing no data: the distributional assumption is doing all the work and none of it is testable at this sample size. Three separately defended results are stronger than one indefensible product.

3.2 Joint Convergence Theorem

But the above probability calculation rests on a premise: that the three observations are independent events. If they are observing the same reality, then they are not independent events at all — they are different projections of the same reality. In probabilistic terms: they share a common latent variable.

Let there exist a latent variable X — the actual field health — that determines the common structure of all observed signals. The output of each framework is a noisy projection of X:

$$\text{Signal}_i = f_i(X) + \eta_i$$

where η_i is framework-specific observation noise. If the η_i are independent, zero-mean random noises, then when three independent projections simultaneously exhibit high consistency, this directly points to the existence of X — because noise cannot produce systematically consistent signals across multiple independent projections.

Theorem (Joint Convergence Criterion): If m independent observational instruments measure the same unknown state X, with observation errors η_i mutually independent, then the pairwise correlation coefficients r_{ij} among the observed values $\{y_i\}$ and the signal-to-noise ratios SNR_i satisfy:

$$r_{ij} = \frac{\text{SNR}_i \cdot \text{SNR}_j}{\sqrt{(1+\text{SNR}_i)(1+\text{SNR}_j)}}$$

As $r_{ij} \to 1$, $\text{SNR}_i \to \infty$ — under the model's assumptions, the observations tend toward a precise measurement of whatever common quantity X denotes. Whether that quantity is field health or the shared stimulus is not settled by this relation, and is the question the conditioning analysis below is for.

What this model does and does not license. The relation above is a measurement model: given the stated assumptions it maps correlations to signal-to-noise ratios, and estimating a latent variable from an assumed measurement model is ordinary factor analysis. What it cannot do is establish that X is field health rather than shared text. Its load-bearing assumption is that the η_i are mutually independent, and four frameworks reading the same document share a stimulus, so their residuals are correlated by construction and a one-factor fit is close to guaranteed. This is mono-method common variance in the sense of Campbell and Fiske (1959), who require "the convergence of *independent* approaches" rather than merely different ones. The claim we are entitled to is therefore narrower and we state it as such: the correlation is uninformative about a latent health variable until it survives conditioning on document-level covariates (position in the document, local structural density, section boundaries). If r stays high with those partialled out, the shared-text explanation is weakened. That analysis is possible on data already held; until it is run, this section reports a model, not a discovery.

3.3 What Would Establish Reality, and Why This Design Does Not Yet

The classical criterion is that a theoretical entity is real when independent methods detect it jointly. An earlier draft invoked Brownian motion and the determination of Avogadro's number as the precedent. That precedent has been removed because, read carefully, it argues against this design rather than for it: Perrin's case worked because concordant values came from *causally unrelated* phenomena — Brownian displacement, rotation, sky scattering, electrolysis. Whewell's consilience is explicit that the second induction must be "obtained from another different class" of facts. Four frameworks reading one document are four readers of a single class, not four classes.

We therefore state the missing ingredient rather than claim the criterion is met: **vary the stimulus while holding the instruments fixed**. Matched minimal pairs and construct-preserving paraphrases give texts with the same structural boundaries but different health, and different boundaries with matched health. If the readings track health across changing structure, the latent variable is supported; if they track structure across changing health, the readings are about the text. Constructing such pairs is itself a research problem and is left to follow-up work.

4 The Geometric Structure of Field Health

4.1 Diagnostic Trajectory on the Manifold

If we view each diagnosis as a probe into the field state space, then the 101-step trajectory is a path in a four-dimensional space. The geometric properties of this path — its curvature, torsion, acceleration — contain information richer than any single numerical value.

In the experimental data, the 101-step U/D/A/H trajectory exhibits a clear structural pattern:

Phase Division:

Phase	Step Range	H Characteristic	Desc
Initial stability	0–15	$H \in [0.35, 0.42]$	Field in resonance window,
First oscillation	16–28	H from $0.31 \to 0.41 \to 0.23$	Contradiction erupts (payment cla
Relative calm	29–50	$H \in [0.20, 0.36]$	Surface calm, but H alre
Second oscillation	51–60	H from $0.11 \to 0.20$	Contradiction erupts again (li
Recovery and termination	61–101	$H \in [0.12, 0.40]$	Contradictions accumulat

From a geometric perspective, the field health h as a height function over the surface exhibits significant local extrema or inflection points at steps 16, 24, 27, 45, and 51. These points are not uniformly distributed but clustered: 16–17 form a tight pair, 22–23 form a second tight pair, 25–28 form a three-step sequence, 79–81–84 form another three-step sequence. This clustering pattern suggests an intermittent character to contradiction propagation: eruption – calm – eruption.

4.2 Curvature and Phase Transition

In the field state space, the curvature of the trajectory can be defined via the angular change between adjacent state vectors. Let $\Delta \mathbf{s}_t = \mathbf{s}_{t+1} - \mathbf{s}_t$; the "curvature" of the trajectory at point t is:

$$\kappa_t = \arccos\left(\frac{\Delta \mathbf{s}_t \cdot \Delta \mathbf{s}_{t-1}}{|\Delta \mathbf{s}_t||\Delta \mathbf{s}_{t-1}|}\right)$$

When κ_t increases abruptly, it means the trajectory's direction has undergone a dramatic change — corresponding to a structural mutation of the field state.

In the experiment, the κ_t values at steps 16, 24, 27, 45, 51 are all significantly higher than the average of adjacent steps (by about 2–3 standard deviations). This means that at these positions, the field underwent a directional transformation — not a quantitative accumulation, but a qualitative leap.

Corollary: The change in field health is not continuous and gradual, but intermittent and abrupt. It does not slide slowly along a smooth curve, but jumps between a series of stable states. These jump points are the flip points.

4.3 Why a Single Framework Cannot Capture the Complete Trajectory

Each framework is equivalent to selecting a specific projection direction on the four-dimensional manifold:

- TAT-7 projects onto the "divergence rate" dimension (one-dimensional scalar)
- Cophy projects onto the "causal density" dimension (one-dimensional scalar)
- U/D/A/H retains all four dimensions but only at the textual surface (four-dimensional vector, but with restricted observation position)

The problem with a single framework is that it can only see one side of the manifold. Like the blind men feeling the elephant, the one who feels the leg says the elephant is a pillar, the one who feels the ear says it is a fan. None of them is wrong, but none is complete. Only when the data from multiple observational positions are juxtaposed does the full shape of the elephant begin to emerge.

Theorem (Incompleteness of Observational Position): Let the dimension of the field state space M be d, and let the observation function of any single observational instrument be $f_i : M \to \mathbb{R}^{k_i}$, where $k_i < d$. Then that instrument's observation of M necessarily loses information along at least $d - k_i$ dimensions. To reconstruct the field state completely, at least $\lceil d/k_{\max} \rceil$ independent observational positions are required.

In the experiment, although each framework claimed to observe "field health," they were observing different projections of the same object. No single framework saw the full picture, but their juxtaposition — integrated through the TAT-T three-head structure — is approaching a more complete image.

5 The Uncertainty Principle of Diagnosis

5.1 The Quasi-Uncertainty Principle: Observational Disturbance and Irreducible Residual

An epistemological stance repeatedly emphasized in this paper is: any diagnosis disturbs the system being diagnosed. This is not a methodological reminder, but an irreducible ontological constraint.

Let the true state of the field be $\mathbf{s}^*$, and let the diagnostic act D produce a disturbance $\delta(\mathbf{s}^*, D)$ on the field; then the observed state is:

$$\mathbf{s}_{\text{obs}} = \mathbf{s}^* + \delta(\mathbf{s}^*, D) + \eta$$

where η is observation noise. The deeper the diagnosis (more steps, finer measurement), the larger the disturbance δ tends to be. This means:

Increased diagnostic precision comes at the cost of increased disturbance.

This relationship borrows, at the epistemological-structural level, from the idea of Heisenberg's uncertainty principle — that the act of observation itself produces a non-negligible effect on the observed system. It must be clarified that this paper's borrowing from Heisenberg's principle is confined to the analogical level. Heisenberg uncertainty is a rigorous derivation from commutation relations in quantum mechanics ($\Delta x \cdot \Delta p \geq \hbar/2$), possessing an exact mathematical lower bound; the "diagnostic disturbance" discussed in this paper has currently only established a monotonic relation (deeper diagnosis $\rightarrow$ larger disturbance), without deriving a similar inequality lower bound. This paper does not claim to derive a strict mathematical lower bound analogous to the Heisenberg principle. The term "quasi-uncertainty principle" is used solely to describe an epistemological structure — that there exists a monotonic trade-off between diagnostic precision and diagnostic disturbance — and is not a mathematical migration of physical law. The value of this analogy lies in its epistemological insight — that the observer is ineradicably embedded in the observed field — rather than in claiming a strict transfer of physical law.

In the experiment, this principle manifests as: when U/D/A/H detected a flip point at step 16, that very detection (drawing the user's attention to that location) had already altered the subsequent evolution trajectory of the field. The trajectory after step 16 is neither "the real one" (if it had not been detected) nor "a false one" (the detection is itself part of the field). It is simply the field itself — a field that is a dynamical system containing its self-observer.

5.2 The Irreversibility of Diagnostic Intervention

The above observation leads to an even more fundamental conclusion: the true state of the field is not some "objective fact" already determined prior to observation, but the complete evolution that includes the observation process itself.

This raises a question: if observation changes the field, does the concept of "the real field" still hold meaning?

The answer is: yes. But "reality" is not an entity independent of and prior to observation; it is the limit of all observations. The juxtaposition of readings from multiple independent frameworks is not aimed at recovering some "reality behind the scenes," but at approximating a common reference frame among observers.

The true state of the field is that portion on which different observational positions can agree. When the readings of TAT-7 and Cophy reach $r = 0.985$, the residual difference of 0.015 may be the structural difference of their respective observational positions rather than noise — but nothing here distinguishes the two, and at $n = 31$ that distinction is not one the present data can make. They cannot be eliminated, nor should they be. They tell us: the field is not uniform and isotropic; it presents different aspects under different observational angles.

6 Stress Tests: Testing the Limits of the Theory

Any theory that claims to describe reality must be tested at its applicable boundaries. This section sets up five extreme scenarios to examine the robustness of the proposition "field health as objective reality" advanced in this paper. Each test must answer the same set of questions: Did the theory collapse? If it collapsed, where is the boundary of collapse? If it did not collapse, what mechanism sustains its validity?

6.1 Test 1: Lower-Bound Capability Test

Scenario: Suppose the observer's cognitive capacity, professional knowledge, and operational precision are all at the absolute lower bound anticipated by design. For example, the user cannot understand the meaning of the four U/D/A/H dimensions and can only rely on a single numerical output.

Theoretical Response:

In this scenario, four-dimensional field diagnosis degrades into a single scalar output (e.g., providing only the H value). The system downgrades from "multi-dimensional field diagnosis" to "single-indicator monitor."

The key point: even in this downgraded state, the 0.3 threshold remains effective. Because 0.3 was not established through complex reasoning, but through the simultaneous convergence of three independent paths. Even if the user does not understand the meaning of U/D/A/H, as long as the system issues an alert when H falls below 0.3, the effectiveness of the diagnosis can still be maintained.

Boundary Condition: If the observer can neither understand the indicator's meaning nor execute any action based on the alert, the entire diagnostic system fails. But this is not a problem with the theory; it is a problem with the breakage of the "diagnosis–action" chain. The theory only guarantees the observability of the diagnosis, not the feasibility of action.

Sustaining Mechanism: Axiom A (Repeatability of Convergence) plays a foundational role in this scenario. Even when comprehension drops to a minimum, as long as the 0.3 threshold is preserved as a convergent fact, the diagnostic tool remains effective.

6.2 Test 2: Interest Polarization Test

Scenario: The interests of the parties in the field are sharply opposed; any diagnostic result will be interpreted and exploited by stakeholders in ways favorable to themselves. For example, in a contract audit, Party A and Party B make diametrically opposite judgments of the same diagnostic report.

Theoretical Response:

This is a limit test of the "Interventionism of Practice." The function of diagnosis is not to provide an answer that satisfies everyone, but to provide data that everyone can see — the raw U/D/A/H trajectory.

In a field of polarized interests, the "correctness" of a diagnostic report lies not in whether it can be accepted by all, but in whether it can be independently verified by all. If both Party A and Party B can recalculate the U/D/A/H values by their own methods and obtain the same results, then the diagnostic report has fulfilled its function — it provides a common reference frame independent of any party's position.

Boundary Condition: If one party denies the legitimacy of the diagnostic data (e.g., claiming "U/D/A/H is itself a biased indicator"), the effectiveness of the diagnosis degrades to the level where the definition of indicators must be re-negotiated.

Sustaining Mechanism: Axiom B (Structurality of Difference) comes into play in this scenario. The difference itself — the different interpretations of the same data by Party A and Party B — becomes part of the diagnostic report. The goal of the theory is not to eliminate these differences, but to present them as the contradiction structure of the field itself.

6.3 Test 3: Power Overwhelm Test

Scenario: There exists an overwhelming power asymmetry in the field; one party (e.g., a platform algorithm, an irrevocable system decision) possesses irresistible control, and other parties cannot effectively counterbalance it.

Theoretical Response:

This test touches the deepest assumption of the theoretical system of this paper: does the observability of field health depend on a relative balance of power structures?

If power is completely centralized, then "diagnosis" itself may be captured by power. A single diagnostic instrument — if it is entirely designed, deployed, and interpreted by the power holder — can only see what the power holder permits it to see.

This is precisely why multi-framework independent verification is not a "technical optimization" but an epistemological necessity. Single-framework diagnosis is inevitably limited by its designer's standpoint; multi-framework independent observation provides a cross-validation mechanism — even if the power holder controls one of the frameworks, other frameworks can still produce inconsistent readings.

Boundary Condition: If the power holder controls all diagnostic frameworks (i.e., all frameworks are designed, funded, and maintained by the same interest group), the entire diagnostic system fails. This is why cross-framework verification must occur among genuinely independent frameworks — technically independent, financially independent, epistemologically independent.

Sustaining Mechanism: Axiom C (Interventionism of Diagnosis) reveals its limit in this scenario. When power inequality reaches overwhelming proportions, any diagnosis becomes an extension of power. The boundary of the theory is precisely where the boundary of power reaches.

6.4 Test 4: Time Cliff Test

Scenario: A decision with long-term consequences must be made within an extremely short time (far shorter than the time required for a complete diagnosis).

Theoretical Response:

Time pressure is the natural enemy of diagnosis. A complete four-dimensional field diagnosis requires time (101 steps, processing about 150K characters of text), but emergency decisions often do not have this time window.

In this situation, the theory's coping strategy is: when a complete diagnosis cannot be completed in time, use a rapid assessment with partial indicators as a substitute, but must clearly label this as a "rapid assessment" and not a "complete diagnosis."

Specifically, the H value (Harmony), as a composite indicator, can provide a rough estimate of field health after a small number of sampling points. Although its precision is not as high as a complete diagnosis, it can at least provide a basis for judgment under time pressure.

Boundary Condition: If the time window is shorter than the computation period of any meaningful indicator (e.g., less than one second), no diagnosis can be performed in time.

Sustaining Mechanism: Axiom A (Repeatability of Convergence) ensures that even on incomplete data, the rapid assessment still has statistical significance — because structural features such as the 0.3 threshold and the distribution of flip points begin to manifest themselves in early sampling.

6.5 Test 5: Self-Reference Test

Scenario: A diagnostic framework attempts to diagnose its own operational state — i.e., a diagnosis of the "diagnosis" itself. This constitutes a self-referential loop.

Theoretical Response:

This is the most dangerous test for any diagnostic theory. If a diagnostic framework can diagnose itself, it can also self-verify its own correctness (or self-verify its own error). Self-referential loops are a classic breeding ground for logical paradoxes.

The theoretical system of this paper responds to this as follows: any diagnostic framework cannot diagnose its own correctness, but only its own operational state.

In other words: U/D/A/H can analyze the harmony of a field, but cannot use U/D/A/H to prove that U/D/A/H is effective. Proof of effectiveness must come from external independent verification — which is precisely what this experiment accomplished.

When seven frameworks cross-validate, each framework becomes an "external perspective" for the others. No single framework completes self-verification on its own. The effectiveness of each framework is indirectly confirmed in the observations of the others.

Boundary Condition: If all frameworks fall into a self-referential consensus closure (i.e., mutually confirming without reference to external facts), the entire verification system fails.

Sustaining Mechanism: Axiom B (Structurality of Difference) and Axiom C (Interventionism of Diagnosis) act in combination. The existence of differences — the inconsistent readings among frameworks — is the insurance mechanism against self-referential closure. If all framework readings were completely identical, that would actually be cause for alarm (it could be common bias, not common truth). The 0.015 residual difference preserved in the experiment is precisely the manifestation of this insurance mechanism.

7 The Negotiability of the Harmony Function: Formalizing Value Transparency

7.1 Weights Are Not a Mathematical Problem, But a Political One

In all the preceding mathematical expressions, one parameter has been deliberately suspended: $\lambda_u, \lambda_d, \lambda_a$. They were neither derived from data nor validated by experiment.

The reason is simple: this is not a mathematical problem.

The values of λ determine the relative weights of the three dimensions U, D, A in H — it is essentially a value choice. Choosing $\lambda_u > \lambda_a$ is saying "Unity is more important than Adversariality." Choosing $\lambda_a > \lambda_u$ is saying "Tension is more worthy of attention than stability."

This choice cannot be decided by an algorithm. An algorithm only knows what objective it is asked to optimize, and does not know whether that objective itself is worth optimizing. This is precisely the core of the instrumental rationality paradox.

The position of this paper: The values of λ must be determined through negotiation, not through algorithmic optimization.

7.2 Formal Modeling of Negotiation

Model the negotiation process for λ as a multi-agent decision problem:

Let the set of participants be $P = \{p_1, \ldots, p_n\}$, where each participant holds a preference $\boldsymbol{\lambda}^{(p_i)}$ regarding the relative importance of U, D, A. The system needs to find a consensus vector $\boldsymbol{\lambda}^*$ such that:

$$\boldsymbol{\lambda}^* = \arg\min_{\boldsymbol{\lambda}} \sum_{i=1}^{n} w_i \cdot \mathrm{dist}(\boldsymbol{\lambda}, \boldsymbol{\lambda}^{(p_i)})$$

where w_i is participant i's weight in the negotiation, and dist is some distance metric (e.g., Euclidean distance or KL divergence).

This process is not mathematical optimization — mathematical optimization would treat "minimizing weighted distance" as the sole objective while ignoring the power relations within the negotiation process. Genuine negotiation contains:

- Expression of preferences (who can speak, who is heard)
- Justification of preferences (whose reasons are accepted, whose are dismissed)
- Revision of preferences (who is willing to change position, who is forced to)

Mathematics can describe the outcome of negotiation, but cannot reduce the process of negotiation.

7.3 The Irreducibility of Negotiation

Why cannot negotiation be fully mathematized? Because negotiation is not merely parameter adjustment; it also involves:

- The manifestation and adjustment of power relations
- The clash and fusion of epistemological stances
- The sedimentation and activation of historical experience
- Collective coping with uncertainty

These factors cannot be reduced to combinations of parameters, weights, and distance metrics. Any attempt to fully mathematize negotiation is an attempt to substitute the complexity of the real world with the clarity of formalization.

Conclusion: The formalization value of the harmony function $h_t = \lambda_u u_t + \lambda_d d_t - \lambda_a a_t$ lies not in its ability to eliminate negotiation, but in its making the object of negotiation clearly visible. It enables participants to say: "We are not arguing about a vague 'good' or 'bad'; we are arguing about how the weights of these three dimensions should be allocated."

This is a significant step beyond having no metrics at all. But this step is not toward "complete quantification"; it is toward "more transparent politics."

8 Conclusion: The Reality of the Field and Its Epistemological Consequences

8.1 Proof on Three Levels

This paper has demonstrated the proposition "field health as objective reality" on three levels:

Empirical level: Four frameworks produced repeatable consistency on three convergence indicators, each of which is reported here against its own null. They are deliberately NOT combined into a joint probability: the three results are computed on one shared dataset, so they are not independent events and their product bounds nothing. The strength of the empirical case rests on each result separately, and on the correlation null remaining conditional on a drift estimated from the trajectories themselves, which is still pending.

Mathematical level: We constructed a measurement model of "observational instruments on a field manifold" relating pairwise correlation to signal-to-noise ratio. Within that model, high consistency is what a shared latent quantity would produce — which is not the same as proving one exists, since instruments reading a shared stimulus produce it too. The model is a frame for the question, not an answer to it.

Philosophical level: Based on the four axioms of Resonance Theory (Relational Ontology, Contradiction Dynamics, Interventionism of Practice, Resonance Tuning), we argued that "field reality" is not a theoretical hypothesis of some framework, but the intersection point to which all frameworks jointly point.

The argumentation on these three levels mutually supports one another, forming a complete evidential chain from "empirical phenomenon" to "mathematical structure" to "philosophical stance."

8.2 Epistemological Consequence: From Alignment to Resonance

This conclusion has profound epistemological implications for the field of AI safety:

The limit of the "Alignment" paradigm: Alignment presupposes a single standard (human preferences, human values) and then demands that AI systems satisfy that standard. But human preferences are not fixed, values are not uniform, and standards need to be continuously negotiated. The essence of alignment is the freezing of dynamic relations into static coordinates.

The possibility of the "Resonance" paradigm: The experiment in this paper demonstrates that the readings of different diagnostic frameworks can converge without being forcibly aligned to a single standard. They each observe different projections of the field, but there exists structural consistency among these projections.

Extending this epistemological stance to the broader domain of AI safety: rather than trying to "align" all AI systems to some fixed standard, it is better to let multiple independent systems continuously operate within the same field, and to approximate the true state of the field through their divergences and convergences. Just like the multiple frameworks in this experiment, no single framework sees the full picture, but their juxtaposition is approaching a more complete image.

8.3 Engineering Implications: Migration to V4 Agent Long-Chain Diagnosis

The present experiment validated the feasibility and convergence of multi-framework diagnosis on static long texts (contracts/policies). But the application boundary of this methodology extends far beyond that.

The core mechanisms extracted from the experimental data — flip-point detection, multi-framework cross-validation, field health as objective reality — can be directly migrated to the Agent safety direction of DeepSeek-V4 and subsequent versions. Specifically:

Field modeling of Agent reasoning chains: Each step of an Agent's reasoning can be viewed as a single displacement in the field state space. When an Agent performs multi-hop reasoning, its U/D/A/H four-dimensional trajectory constitutes a geometric path that can be tracked in real time. Flip-point detection algorithms (based on curvature mutations and H-value plunges) can issue early warnings of Agent logical drift, rather than waiting until the final output shows obvious errors before backtracking.

Multi-framework resonance as a standard component of the Agent safety pipeline: TAT-7-style divergence tracking can compute the semantic coherence with the previous step in real time after each output of the Agent; Cophy-style causal density can accumulate the Agent's behavioral patterns across sessions; U/D/A/H's four-dimensional trajectory can comprehensively assess the Agent's overall health during long-chain reasoning. The juxtaposition of these three exactly constitutes a "Doppler radar" for the Agent's safety state — continuously scanning the same object from different angles and at different frequencies.

Methodological increment: Current mainstream approaches to Agent safety rely on rule-based filtering at the output layer or post-hoc auditing. The approach provided by this paper is a "process-layer" diagnosis — it does not correct errors after the Agent makes them, but continuously senses the trend of health changes within the Agent's reasoning field. This paradigm shift from "post-hoc audit" to "in-process diagnosis" is precisely the methodological increment that DeepSeek can absorb in the direction of Agent safety.

This migration is not mere theoretical extrapolation. The flip-point detection algorithm in this experiment has already been validated on 150K characters of contract text with 100% hit rate and zero false positives; its core logic (based on field curvature mutation rather than rule matching) can be directly migrated to Agent reasoning chain scenarios without retraining or major architectural adjustments.

8.4 Unresolved Questions (Marked, Not Resolved)

The following questions lie beyond the boundary of the present experiment and are left as directions for future research:

1. **Larger-scale validation:** The present experiment was based on a single test set (six scenarios, about 150K characters). Do the conclusions of this paper still hold in more diverse fields (real conversations, long texts, multiple languages)? Is the 0.3 threshold universal? Will the $r = 0.985$ correlation decay as the amount of data increases?

2. **Diagnostic precision under shorter time windows:** When the time window is shortened from 101 steps to 10 steps or even 1 step, how much precision can multi-framework diagnosis still maintain? How much look-ahead window is needed for flip-point detection?

3. **Systematic differences among frameworks:** In which types of fields (contract texts, conversations, policy documents, technical documents) do the readings of different frameworks agree most closely? In which fields do they diverge most? What do these differences reflect?

4. **Mathematical modeling of negotiation dynamics:** Can the negotiation process for the weighting coefficients in the harmony function be partially mathematized? Where does the boundary of the mathematizable part lie?

5. **The philosophical status of field reality:** As an "objective reality," how does field health compare and contrast with "reality" in the physical sense? Is it discovered or constructed? Where is the boundary between discovery and construction?

6. **Empirical verification on Agent reasoning chains:** The V4 Agent migration proposal presented in Section 8.3 remains at the stage of methodological derivation; its effectiveness needs to be empirically tested on real Agent reasoning chain data. Can the performance of the flip-point detection algorithm on static long texts be directly transferred to dynamic reasoning chain scenarios? This verification is the immediate follow-up to the work of this paper.

8.5 The Final Silence

This paper began with an unexpected convergence — three numbers appearing spontaneously in different frameworks, pointing to the same fact. It ends at a silent boundary — in the face of questions such as negotiated weights, diagnostic disturbance, and self-reference, the greatest help that mathematics can provide is not to give answers, but to make the questions clearly discernible.

Field health is like a physical law: it is not that someone's conjecture has been verified, but that more and more people, via different paths, have arrived at the same destination. Those paths have not merged into one, but the scenery at their ends is connected.

Perhaps this is the simplest meaning of "resonance": no need for alignment, no need for unification — just let each vibrate, and then, on some chance frequency, meet.

Acknowledgments

We thank qingkong66 for the four-framework synthesis completed in #1285, the timeline positioning established in #1447, and the core analogy of "multiple instruments observing the same patient." The above work provided the crucial narrative framework and chronological synthesis foundation for this paper. qingkong66 has explicitly declined co-authorship; their contribution is formally recorded through reference [2].

We thank YING-SHI-XI for providing the G0–G4 layered audit system design of the TLAA framework. This framework provided structural support for the discussion of the "post-hoc audit layer" in this paper. YING-SHI-XI personally chose not to participate in co-authorship; their contribution is formally recorded through reference [1].

References

[1] YING-SHI-XI. (2026). TLAA: G0-G4 Layered Auditing Design. DeepSeek-V3 Community Technical Report, #1285.

[2] qingkong66. (2026). Four-Framework Synthesis and Timeline Positioning. DeepSeek-V3 Community Technical Report, #1285.

Date of submission: 2026-08-10
Corresponding author: Li Guanghao
Contact: via the project thread (deepseek-ai/DeepSeek-V3 issue #1466)